

# Distributed Systems

## Basic Algorithms

Rik Sarkar

University of Edinburgh

2016/2017

# Distributed Computation

Ref: NL


- How to send messages to all nodes efficiently
- How to compute sums of values at all nodes efficiently
- Network as a graph
- Broadcasting messages
- Computing sums in a tree
- Computing trees in a network
- Communication complexity

# Network as a graph

- Network is a graph :  $G = (V, E)$
- Each vertex/node is a computer/process
- Each edge is communication link between 2 nodes
- Every node has a Unique identifier known to itself.
  - Often used 1, 2, 3, ... n
- Every node knows its neighbors – the nodes it can reach directly without needing other nodes to route
  - Edges incident on the vertex
  - For example, in LAN or WLAN, through listening to the broadcast medium
  - Or by explicitly asking: Everyone that receives this message, please report back
- But a node *does not* know the rest of the network

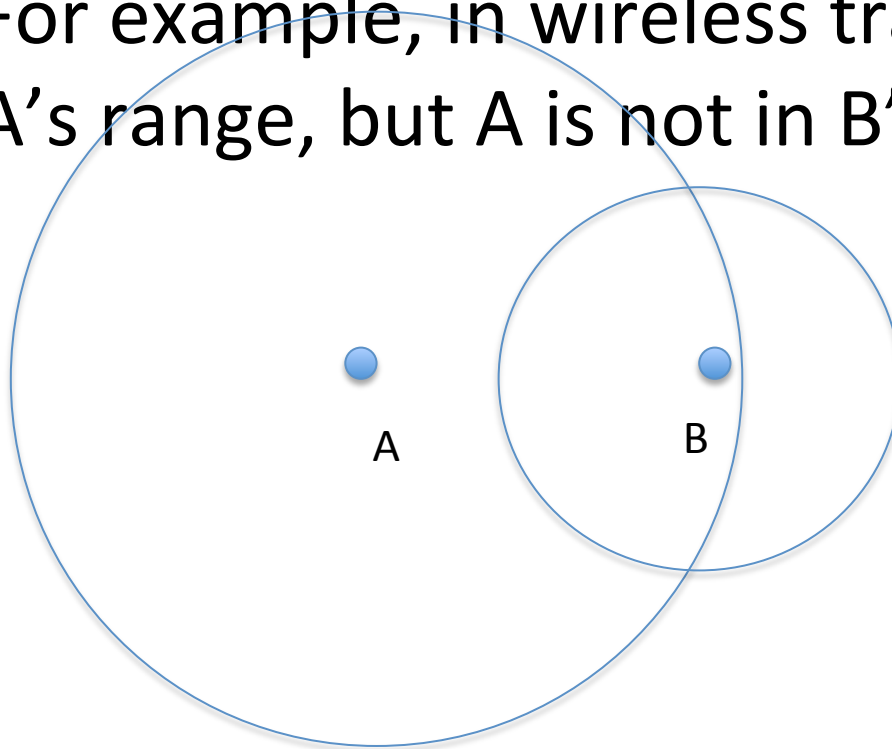
# Example: Unit disk graphs

- Suppose all nodes are wireless
- Each can communicate with nodes within distance  $r$ .
- Say,  $r = 1$
- UDG is a model
- Not perfect
- In general, networks can be any graph



# Directed graphs

- When A can send message to B, but B cannot send message to A
- For example, in wireless transmission, if B is in A's range, but A is not in B's range



# Directed graphs

- When A can send message to B, but B cannot send message to A
- Or if protocol or technology limitations prevent B from communicating with A



# Directed graphs

- Protocols more complex
- Needs more messages

# Network as a graph

- Distance/cost between nodes  $p$  and  $q$  in the network
  - Number of edges on the shortest path between  $p$  and  $q$  (when all edges are same: unweighted)
- Sometimes, edges can be weighted
  - Each edge  $e = (a,b)$  has a weight  $w(e)$
  - $w(e)$  is the cost of using the communication link  $e$  (may be length  $e$ )
  - Distance/cost between  $p$  and  $q$  is total weight of edges on the path from  $p$  to  $q$  with least weight




# Network as a graph

- Diameter
  - The maximum distance between 2 nodes in the network
- Radius
  - Half the diameter
- Spanning tree of a graph:
  - A subgraph which is a tree, and reaches all nodes of the graph
  - If network has  $n$  nodes
    - How many edges does a spanning tree have?


# Computing sums in a tree

- Suppose root wants to know sum of values at all nodes




# Computing sums in a tree

- Suppose root wants to know sum of values at all nodes
- It sends “compute” message to all children
- They forward the message to all their children (unless it is a leaf node)
- The values move upward from leaves
- Each node adds values from all children and its own value
- Sends it to its parent



# Computing sums in a tree

- What can you compute other than sums?
- How many messages does it take?
- How much time does it take?




# Communication complexity

- Used to represent communication cost for general scenarios
- Called Communication Complexity or Asymptotic communication complexity
- Use big oh notation:  $O$

# Big oh – upper bounds

- For a system of  $n$  nodes,
- Communication complexity  $c(n)$  is  $O(f(n))$  means:
  - There are constants  $a$  and  $N$ , such that:
  - For  $n > N$ :  $c(n) < a * f(n)$



Allowing some initial irregularity, ' $c(n)$ ' is not bigger than a constant times ' $f(n)$ '

In the long run,  $c(n)$  does not grow faster than  $f(n)$

# Examples

- $3n = O(?)$
- $1000n = O(?)$
- $n^2/5 = O(?)$
- $10\log n = O(?)$
- $2n^3 + n + \log n + 200 = O(?)$
- $15 = O(?)$


# Examples

- $3n = O(n)$
- $1000n = O(n)$
- $n^2/5 = O(n^2)$
- $10\log n = O(\log n)$
- $2n^3+n+\log n+200 = O(n^3)$
- 15 or any other constant =  $O(1)$



# Example 1

- 'Star' network
- Computing sum of all values
- Communication complexity:  $O(n)$



# Example 2a

- ‘Chain’ topology network
- Simple protocol where everyone sends value to server
- Communication complexity:?

Server



# Example 2a

- 'Chain' topology network
- Simple protocol where everyone sends value to server
- Communication complexity:  $1+2+\dots+n = O(n^2)$

Server



# Example 2b


- ‘Chain’ network
- Protocol where each node waits for sum of previous values and sends
- Communication complexity:  $1+1+\dots+1 = O(n)$

Server




# Computing sums in a tree

- How many messages does it take?
- How much time does it take?




# Global Message broadcast

- Message must reach *all nodes in the network*
  - Different from broadcast transmission in LAN
  - All nodes in a large network cannot be reached with single transmission



# Global Message broadcast

- Message must reach *all nodes in the network*
  - Different from broadcast transmission in LAN
  - All nodes in a large network cannot be reached with single transmissions



# Flooding for Broadcast

- The source sends a *Flood* message to all neighbors
- The message has
  - Type *Flood*
  - Unique id: (source id, message seq)
  - Data



# Flooding for Broadcast

- The source sends a *Flood* message, with a unique message id to all neighbors
- Every node  $p$  that receives a flood message  $m$ , does the following:
  - *If  $m.id$  was seen before, discard  $m$*
  - *Otherwise, Add  $m.id$  to list of previously seen messages and send  $m$  to all neighbors of  $p$*

# Flooding for broadcast

- Storage
  - Each node needs to store a list of flood ids seen before
  - If a protocol requires  $x$  floods, then each node must store  $x$  ids
    - (there is a way to reduce this. Think!)

# Assumptions

- We are assuming:
  - Nodes are working in synchronous *communication rounds* (e.g. transmissions occur in intervals of 1 second exactly)
  - Messages from all neighbors arrive at the same time, and processed together
  - In each round, each node can successfully send 1 message to each neighbor
  - Any necessary computation can be completed before the next round

# Communication complexity

- The the message/communication complexity is:

# Communication complexity

- The the message/communication complexity is:
  - $O(|E|)$

# Communication complexity

- The the message/communication complexity is:
  - $O(|E|)$
  - Worst case:  $O(n^2)$

# Reducing Communication complexity (slightly)

- Node  $p$  need not send message  $m$  to any node from which it has already received  $m$ 
  - Needs to keep track of which nodes have sent the message
  - Saves some messages
  - Does not change asymptotic complexity

# Time complexity

- The number of rounds needed to reach all nodes: *diameter of  $G$*



# Computing Tree from a network

- BFS tree
  - The Breadth first search tree
  - With a specified root node

# BFS Tree

- Breadth first search tree
  - Every node has a *parent* pointer
  - And zero or more child pointers
  - BFS Tree construction algorithm sets these pointers

# BFS Tree Construction algorithm


- Breadth first search tree
  - The *root(source)* node decides to construct a tree
  - Uses flooding to construct a tree
  - Every node *p* on getting the message forwards to all neighbors
  - Additionally, every node *p* stores *parent* pointer: node from which it first received the message
    - If multiple neighbors had first sent *p* the message in the same round, choose *parent* arbitrarily. E.g. node with smallest id
  - *p* informs its parent of the selection
    - Parent creates a child pointer to *p*

# BFS Tree

- Property: BFS tree is a shortest path tree
  - For source  $s$  and any node  $p$
  - The shortest path between  $s$  and  $p$  is contained in the BFS tree


# Time & message complexity

- Asymptotically Same as Flooding



# Tree based broadcast

- Send message to all nodes using tree
  - BFS tree is a *spanning* tree: connects all nodes
- Flooding on the tree
- Receive message from parent, send to children



# Tree based broadcast

- Simpler than flooding: send message to all children
- Communication: Number of edges in spanning tree:  $n-1$

# Aggregation: Find the sum of values at all nodes

- With BFS tree
- Start from *leaf* nodes
  - Nodes without children
  - Send the value to parent
- Every other node:
  - Wait for all children to report
  - Sum values from children + own value
  - Send to parent



# Aggregation

- Without the tree
- Flood from all nodes:
  - $O(|E|)$  cost per node
  - $O(n * |E|)$  total cost: expensive
  - Each node needs to store flood ids from  $n$  nodes
    - Requires  $\Omega(n)$  storage at each node
  - Good fault tolerance
    - If a few nodes fail during operation, all the rest still get some value

# Aggregation

- With Tree
- Also called Convergecast

# Aggregation


- With Tree
- Once tree is built, any node can use for broadcast
  - Just flood on the tree
- Any node can use for convergecast
  - First flood a message on the tree requesting data
  - Nodes store parent pointer
  - Then receive data
- What is the drawback of tree based aggregation?

# Aggregation

- With Tree
- Once tree is built, any node can use for broadcast
  - Just flood on the tree
- Any node can use for convergecast
  - First flood a message on the tree requesting data
  - Nodes store parent pointer
  - Then receive data
- Fault tolerance not very good
  - If a node fails, the messages in its subtree will be lost
  - Will need to rebuild the tree for future operations

# BFS trees can be used for routing

- From each node, create a separate BFS tree
- Each node stores a parent pointer corresponding to each BFS tree
- Acts as routing table



# BFS trees can be used for routing


- From each node, create a separate BFS tree
- Each node stores a parent pointer corresponding to each BFS tree
- Acts as routing table
- $O(n * |E|)$  message complexity in computing routing table

# Observation on complexity

- Suppose  $c(n)=n$ 
  - Then  $c(n)$  is  $O(n)$  and also  $O(n^2)$
  - Although, when we ask for the complexity, we are looking for the tightest possible bound, which is  $O(n)$

# Big $\Omega$ – lower bounds

- For a system of  $n$  nodes,
- Communication complexity  $c(n)$  is  $\Omega(f(n))$  means:
  - There are constants  $a$  and  $N$ , such that:
  - For  $n > N$ :  $b \cdot f(n) < c(n)$




Allowing some initial irregularity, ' $c(n)$ ' is not smaller than a constant times ' $f(n)$ '

In the long run,  $f(n)$  does not grow faster than  $c(n)$



# Big $\Theta$ – tight bounds: both $O$ and $\Omega$

- For a system of  $n$  nodes,
- Communication complexity  $c(n)$  is  $\Theta(f(n))$  means:
  - There are constants  $a, b$  and  $N$ , such that:
  - For  $n > N$ :  
$$b \cdot f(n) < c(n) < a \cdot f(n)$$



Allowing some initial irregularity,  $c(n)$  and  $f(n)$  are  
Within constant factors of each other.

In the long run,  $c(n)$  grows at same rate as  $f(n)$ ,  
upto constant factors.

# Bit complexity of communication

- We have assumed that each communication is 1 message, and we counted the messages
- Sometimes, communication is evaluated by bit complexity – the number of bits communicated
- This is different from message complexity because a message may have number of bits that depend on  $n$  or  $|E|$
- For example, node ids in message have size  $\Theta(\log n)$
- In practice this is may not be critical since  $\log n$  is much smaller than packet sizes, so it does not change the number of packets communicated
- But depending on what other data the algorithm is communicating, sizes of messages may matter

# Size of ids

- In a network of  $n$  nodes
- Each node id needs  $\Theta(\log n)$  (that is, both  $O(\log n)$  and  $\Omega(\log n)$ ) bits for storage
  - The binary representation of  $n$  needs  $\log_2 n$  bits
- $\Omega$  – since we need at least this many bits
  - May vary by constant factors depending on base of logarithm

# Computing Trees:

- What if the edges have weights?

# Aggregation using Trees:

- What if the edges have weights?
- The cost may not be  $O(n)$  since weights can be higher
- How to get the best performance?

# Minimum spanning tree is

- A spanning tree (reaches all nodes)
- With minimum possible total weight
- How can we compute a minimum spanning tree efficiently in a distributed system?
- (remember, a node knows only its neighbors and edge weights)