



Extreme Computing

Hadoop MapReduce in more detail



How will I actually learn Hadoop?

- This class session
- Hadoop: The Definitive Guide
- RTFM
- There is a lot of material out there
 - There is also a lot of useless material
 - You need to filter it
 - Just because some random guy wrote a blog post about something does not make it right
 - We will have lab sessions
 - Attend them
 - Ask questions
 - Ask more questions



Basic Hadoop API

- Mapper
 - `void setup(Mapper.Context context)`
Called once at the beginning of the task
 - `void map(K key, V value, Mapper.Context context)`
Called once for each key/value pair in the input split
 - `void cleanup(Mapper.Context context)`
Called once at the end of the task
- Reducer/Combiner
 - `void setup(Reducer.Context context)`
Called once at the start of the task
 - `void reduce(K key, Iterable<V> values, Reducer.Context context)`
Called once for each key
 - `void cleanup(Reducer.Context context)`
Called once at the end of the task

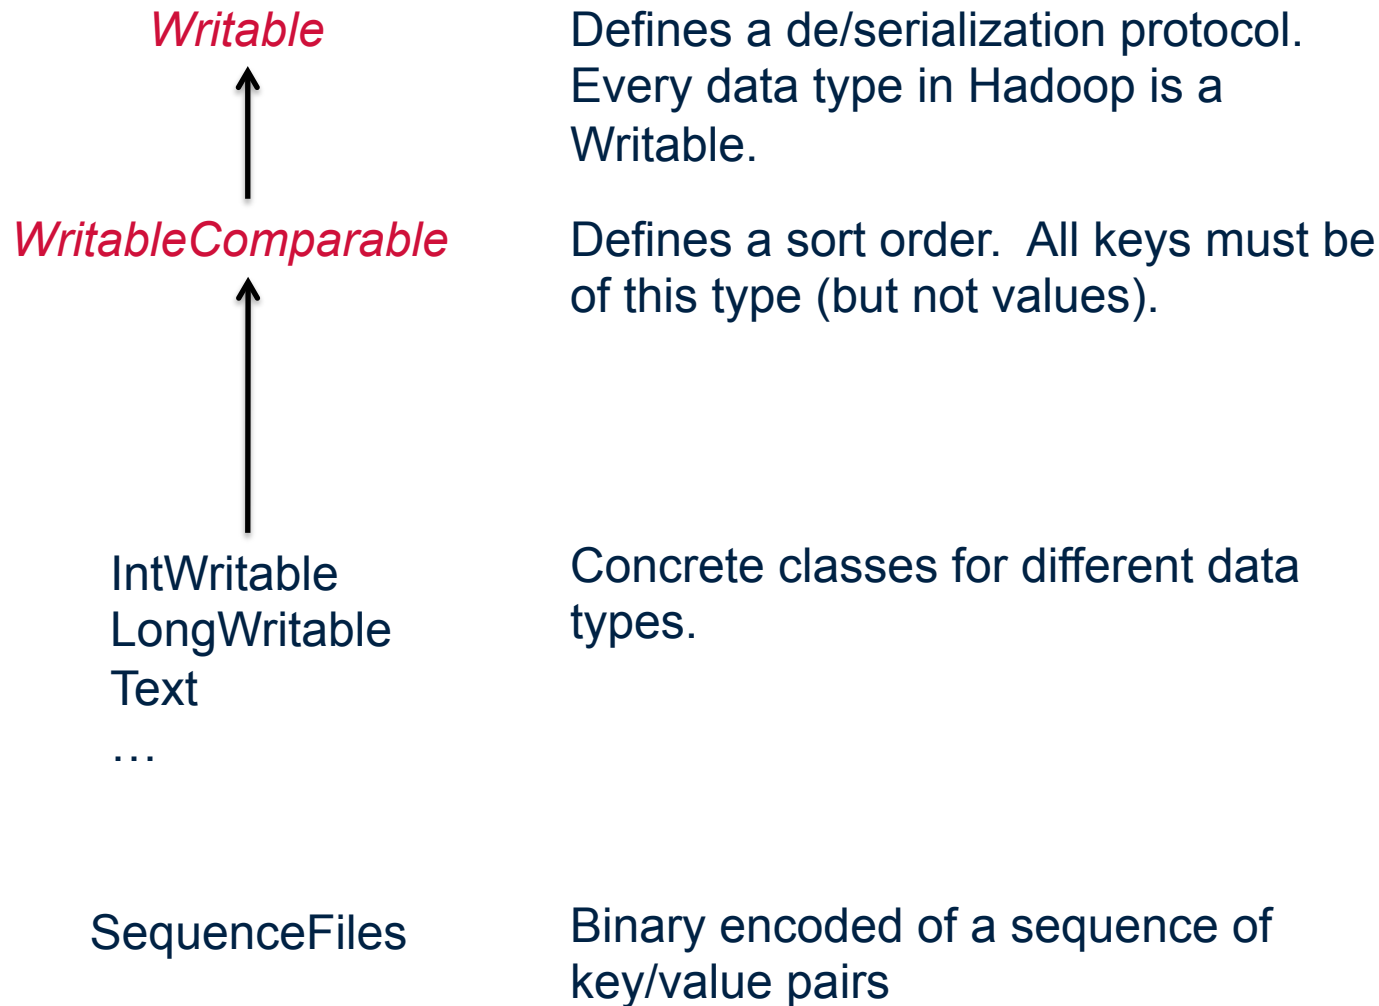


Basic Hadoop API

- Partitioner
 - `int getPartition(K key, V value, int numPartitions)`
Get the partition number given total number of partitions
- Job
 - Represents a packaged Hadoop job for submission to cluster
 - Need to specify input and output paths
 - Need to specify input and output formats
 - Need to specify mapper, reducer, combiner, partitioner classes
 - Need to specify intermediate/final key/value classes
 - Need to specify number of reducers (but not mappers, why?)
 - Don't depend of defaults!



Data types in Hadoop: keys and values





“Hello World”: Word count

```
Map(String docid, String text):  
    for each word w in text:  
        Emit(w, 1);
```

```
Reduce(String term, Iterator<Int> values):  
    int sum = 0;  
    for each v in values:  
        sum += v;  
        Emit(term, value);
```



“Hello World”: Word Count

```
private static class MyMapper extends
    Mapper<LongWritable, Text, Text, IntWritable> {

    private final static IntWritable ONE = new IntWritable(1);
    private final static Text WORD = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = ((Text) value).toString();
        StringTokenizer itr = new StringTokenizer(line);
        while (itr.hasMoreTokens()) {
            WORD.set(itr.nextToken());
            context.write(WORD, ONE);
        }
    }
}
```



“Hello World”: Word Count

```
private static class MyReducer extends
    Reducer<Text, IntWritable, Text, IntWritable> {

    private final static IntWritable SUM = new IntWritable();

    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        Iterator<IntWritable> iter = values.iterator();
        int sum = 0;
        while (iter.hasNext()) {
            sum += iter.next().get();
        }
        SUM.set(sum);
        context.write(key, SUM);
    }
}
```




Getting data to mappers and reducers

- Configuration parameters
 - Directly in the Job object for parameters
- Side data
 - DistributedCache
 - Mappers/reducers read from HDFS in setup method
- Avoid object creation at all costs
 - Reuse Writable objects, change the payload
- Execution framework reuses value object in reducer
- Passing parameters via class statics



Complex Data Types in Hadoop

- How do you implement complex data types?
- The easiest way:
 - Encoded it as Text, e.g., $(a, b) = "a:b"$
 - Use regular expressions to parse and extract data
 - Works, but pretty hack-ish
- The hard way:
 - Define a custom implementation of Writable(Comparable)
 - **Must implement: readFields, write, (compareTo)**
 - Computationally efficient, but slow for rapid prototyping
 - Implement WritableComparator hook for performance
- Somewhere in the middle:
 - Some frameworks offers JSON support and lots of useful Hadoop types

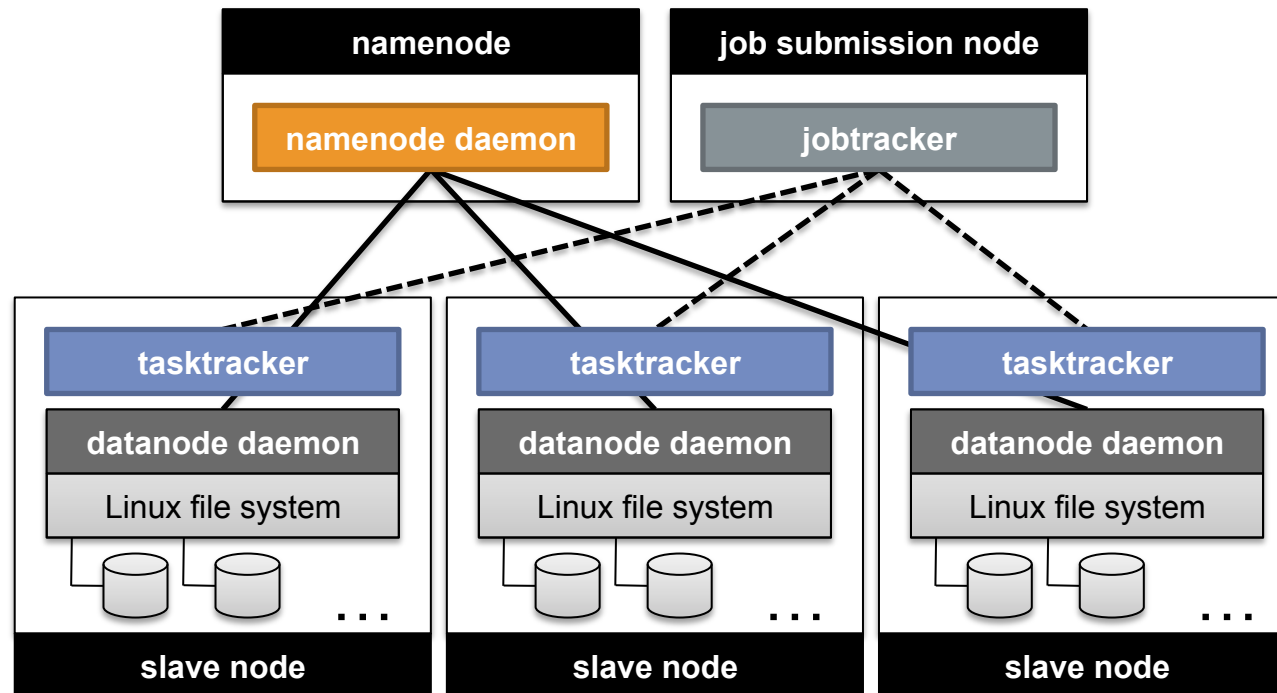


Basic cluster components

- One of each:
 - Namenode (NN): master node for HDFS
 - Jobtracker (JT): master node for job submission
- Set of each per slave machine:
 - Tasktracker (TT): contains multiple task slots
 - Datanode (DN): serves HDFS data blocks



Recap





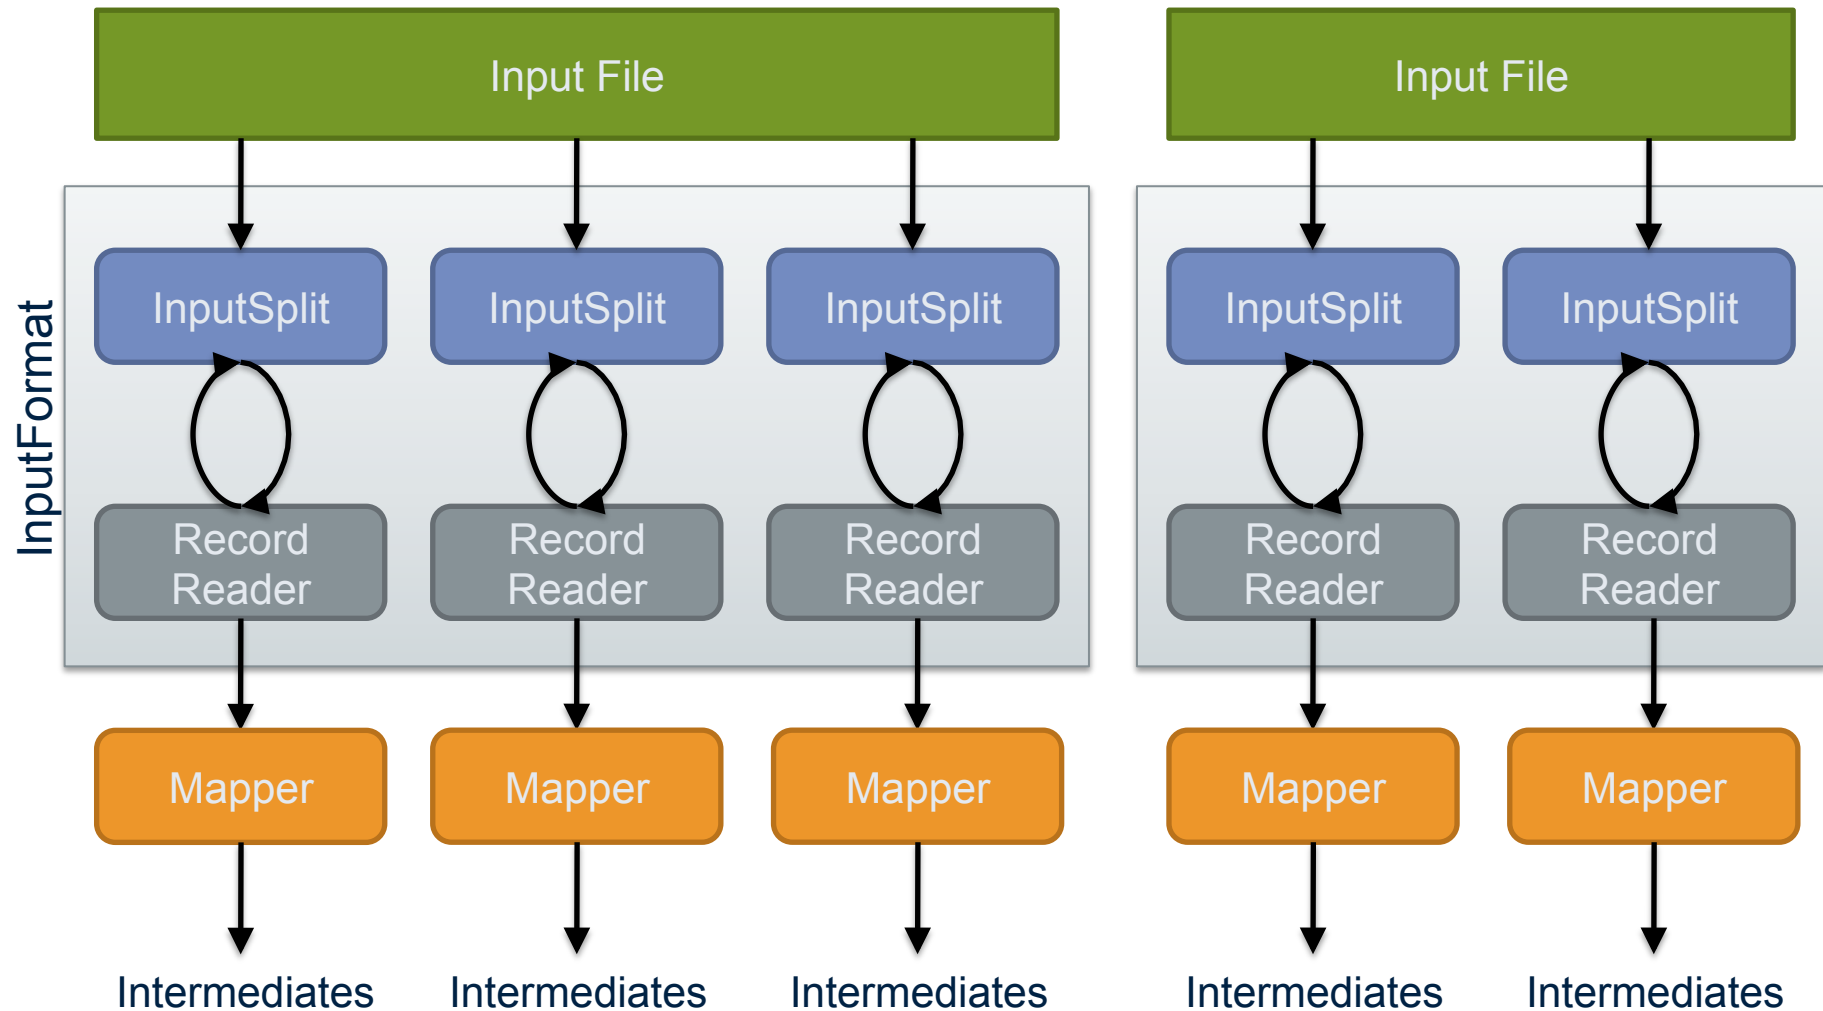
Anatomy of a job

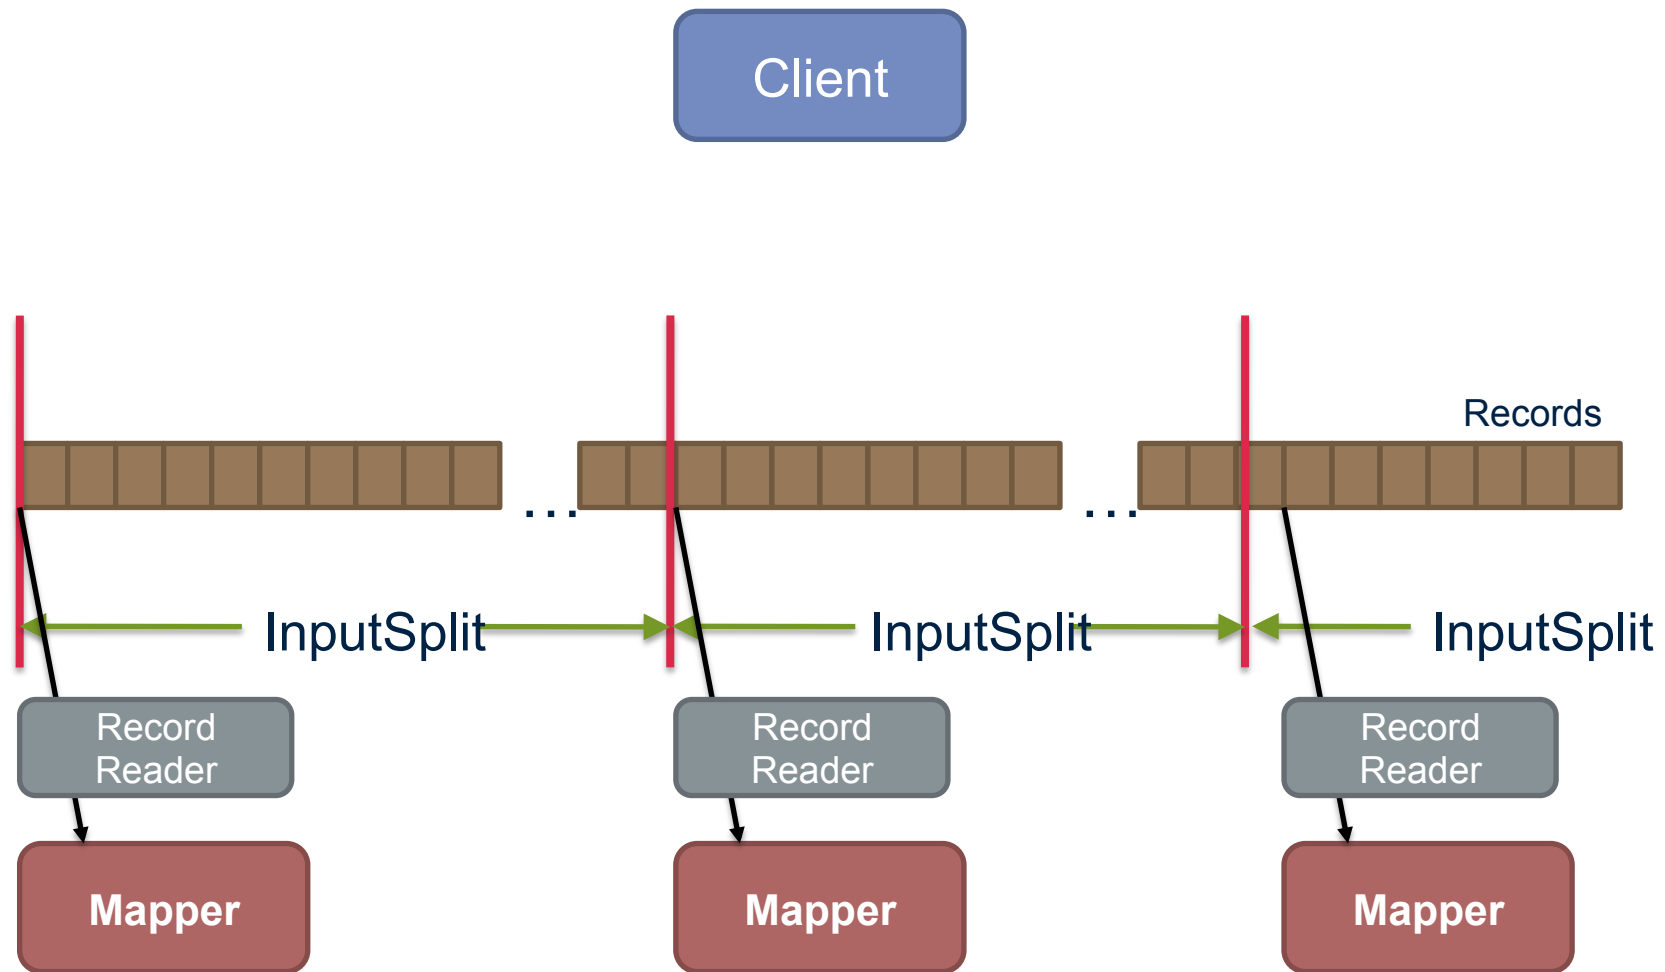
- MapReduce program in Hadoop = Hadoop job
 - Jobs are divided into map and reduce tasks
 - An instance of running a task is called a task attempt (occupies a slot)
 - Multiple jobs can be composed into a workflow
- Job submission:
 - Client (i.e., driver program) creates a job, configures it, and submits it to jobtracker
 - That's it! The Hadoop cluster takes over

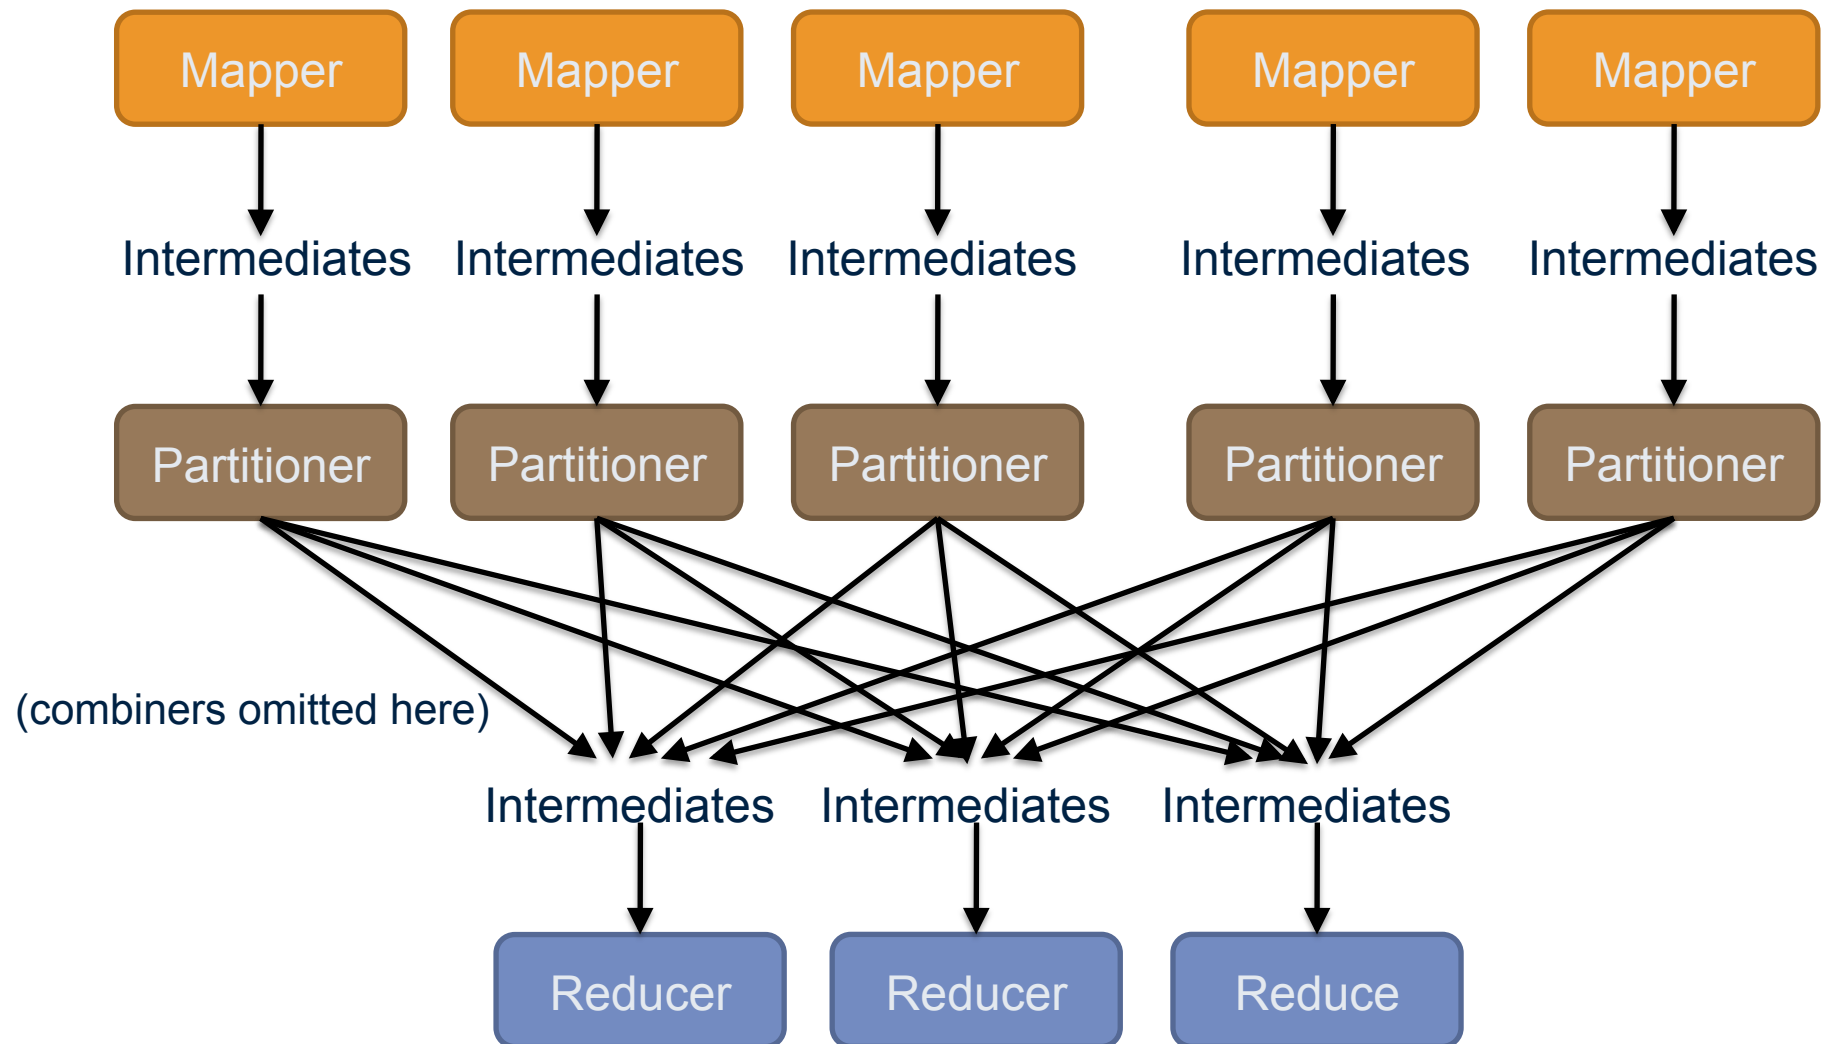


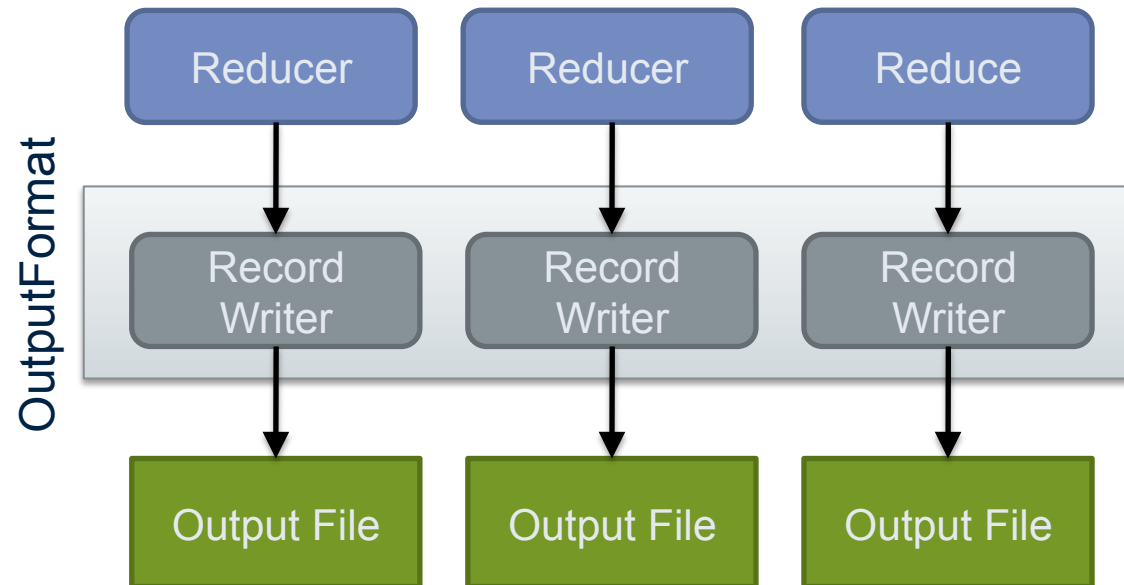
Anatomy of a job

- Behind the scenes:
 - Input splits are computed (on client end)
 - Job data (jar, configuration XML) are sent to JobTracker
 - JobTracker puts job data in shared location, enqueues tasks
 - TaskTrackers poll for tasks
 - Off to the races











Input and output

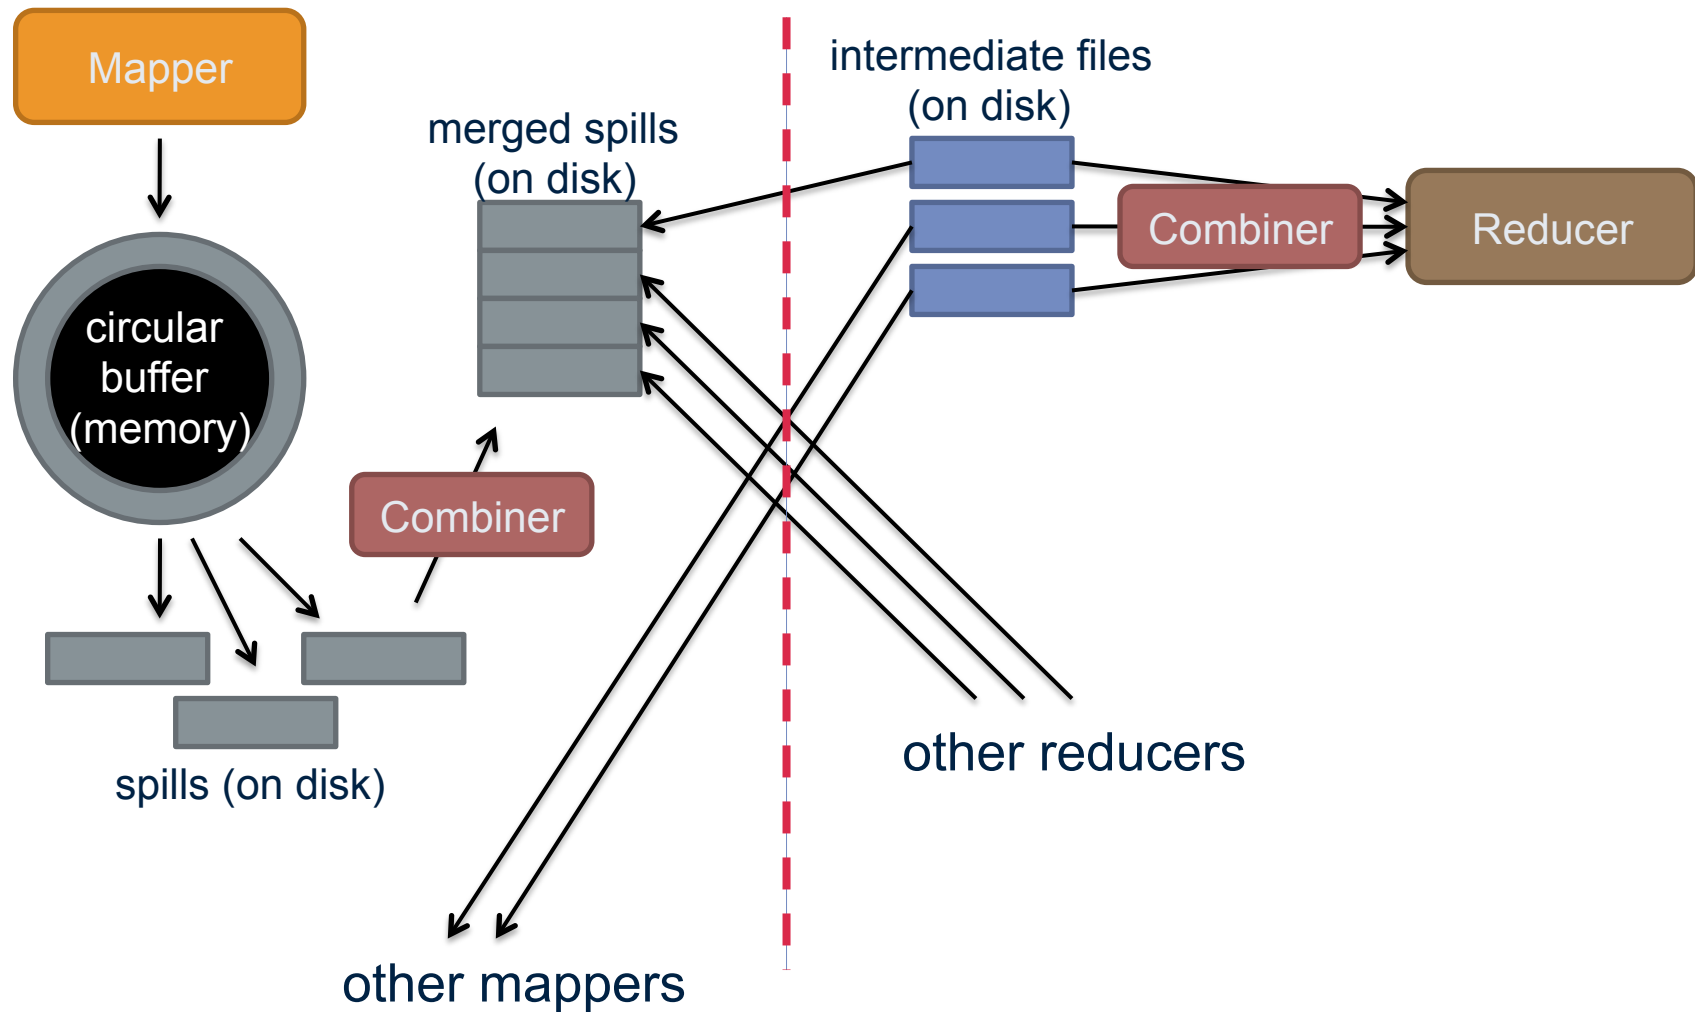
- InputFormat:
 - TextInputFormat
 - KeyValueTextInputFormat
 - SequenceFileInputFormat
 - ...
- OutputFormat:
 - TextOutputFormat
 - SequenceFileOutputFormat
 - ...



Shuffle and sort in Hadoop

- Probably the most complex aspect of MapReduce
- Map side
 - Map outputs are buffered in memory in a circular buffer
 - When buffer reaches threshold, contents are spilled to disk
 - Spills merged in a single, partitioned file (sorted within each partition): combiner runs during the merges
- Reduce side
 - First, map outputs are copied over to reducer machine
 - Sort is a multi-pass merge of map outputs (happens in memory and on disk): combiner runs during the merges
 - Final merge pass goes directly into reducer

Shuffle and sort





Recommended workflow

- Here's one way to work
 - Develop code in your favourite IDE on host machine
 - Build distribution on host machine
 - Check out copy of code on VM
 - Copy (i.e., scp) jars over to VM (in same directory structure)
 - Run job on VM
 - Iterate
- Avoid using the UI of the VM
 - Directly ssh into the VM
- Deploying the job
- `$HADOOP_CLASSPATH`
- `hadoop jar MYJAR.jar -D k1=v1 ... -libjars foo.jar,bar.jar
my.class.to.run arg1 arg2 arg3 ...`



Actually running the job

- \$HADOOP_CLASSPATH
- `hadoop jar MYJAR.jar`
 - `-D k1=v1 ...`
 - `-libjars foo.jar,bar.jar`
 - `my.class.to.run arg1 arg2 arg3 ...`



Debugging Hadoop

- First, take a deep breath
- Start small, start locally
- Build incrementally
- Different ways to run code:
 - Plain Java
 - Local (standalone) mode
 - Pseudo-distributed mode
 - Fully-distributed mode
- Learn what's good for what

We will start by using Python bindings before deploying Java jobs



Hadoop debugging strategies

- Good ol' `System.out.println`
 - Learn to use the webapp to access logs
 - Logging preferred over `System.out.println`
 - Be careful how much you log!
- Fail on success
 - Throw `RuntimeExceptions` and capture state
- Programming is still programming
 - Use Hadoop as the glue
 - Implement core functionality outside mappers and reducers
 - Independently test (e.g., unit testing)
 - Compose (tested) components in mappers and reducers



Summary

- Presented Hadoop in more detail
- Described the implementation of the various components
- Described the workflow of building and deploying applications
- Things are a lot more complicated than this
- We will next turn to algorithmic design for MapReduce