

# Extreme Computing Revision

# Emphasis on solving problems

Design a system to store YouTube videos.

You're Twitter. Make a low-latency feed system.

Show the most frequently visited pages. Live.

# Reason About a Problem

- What's the bottleneck? Compute, storage, bandwidth, latency, ...
- Which data is big? Which is small? → Join strategies
- Is it balanced?
- What if that machine fails?

# Reason About a Problem

- What's the bottleneck? Compute, storage, bandwidth, latency, ...
- Which data is big? Which is small? → Join strategies
- Is it balanced?
- What if that machine fails?

More fun: how will this system fail?

# Bag of Tricks

- Sharding aka partitioning: divide the work across machines
- Replication for speed and fault tolerance
- “Cold” large read-only store, “hot” small mutable store
- Approximations: Bloom filters, streaming counts, resevoir sampling

# Systems

**MapReduce** Parallel batch processing

**BitTorrent** File sharing

**HDFS** File storage

**Chord** Divide responsibility as machines join and leave

**SSTable** Large read-only key-value store

**BigTable** Large mutable key-value store

# Systems

MapReduce Parallel batch processing

BitTorrent File sharing

HDFS File storage

Chord Divide responsibility as machines join and leave

SSTable Large read-only key-value store

BigTable Large mutable key-value store

Given a problem, name a system and apply it.  
Take inspiration from the design of these systems.

# Linearisable versus Sequential

Alice and Bob write checks to each other for the same amount.

## Alice's Statement

|     |                   |
|-----|-------------------|
| -10 | Check Alice → Bob |
| 0   | Check Bob → Alice |

## Bob's Statement

|     |                   |
|-----|-------------------|
| -10 | Check Bob → Alice |
| 0   | Check Alice → Bob |

Both overdraft.

- ✓ Sequential: each client sees a consistent order
- ✗ Linearizable: no globally linear story